

LAB 03

Parallel Programming on PULP

Alberto Dequino (alberto.dequino@unibo.it)

Călin Diaconu (calin.diaconu2@unibo.it)


Francesco Conti

Giuseppe Tagliavini





ALMA MATER STUDIORUM
UNIVERSITÀ DI BOLOGNA



PULP: A Parallel Ultra-Low-Power Computing Platform



PULP: A Parallel Ultra-Low-Power Computing Platform





PULP: A Parallel Ultra-Low-Power Computing Platform



Target a **Shared-Memory** parallel programming model: 4—16 DSP cores accessing a **Shared L1 Memory** (*Tightly Coupled Data Mem* or *TCDM*)
How is memory **organized** to allow this?

PULP: A Parallel Ultra-Low-Power Computing Platform





Target a **Shared-Memory** parallel programming model: 4—16 DSP cores accessing a **Shared L1 Memory** (*Tightly Coupled Data Mem* or *TCDM*)

How is memory **organized** to allow this?



The memory is organized in **Multiple Banks** → **concurrent access**.

PULP: A Parallel Ultra-Low-Power Computing Platform





Target a **Shared-Memory** parallel programming model: 4—16 DSP cores accessing a **Shared L1 Memory (*Tightly Coupled Data Mem* or *TCDM*)**
How is memory **organized** to allow this?
The memory is organized in **Multiple Banks** → **concurrent access**.
Solution: **full connectivity** with 1-cycle latency **crossbar**

PULP: A Parallel Ultra-Low-Power Computing Platform





Multiple separate instruction streams (hardware threads):
MIMD (Multiple Instruction Multiple Data) execution scheme...

PULP: A Parallel Ultra-Low-Power Computing Platform





Multiple separate instruction streams (hardware threads):
MIMD (Multiple Instruction Multiple Data) execution
 scheme... but optimized for a single parallel program (**SPMD**)

PULP: A Parallel Ultra-Low-Power Computing Platform




High-speed synchronization to minimize overhead in **synchronization primitives**

PULP: A Parallel Ultra-Low-Power Computing Platform




PULP Execution Model

The Cluster is **clock-gated** at boot: a single thread runs on the **Fabric Controller** (MCU core)




PULP Execution Model

Activate the cluster and **call** a function on the **cluster core 0**




PULP Execution Model

Fork an **execution team** on multiple cluster cores; they can be synchronized with *barriers* and *critical sections*




PULP Execution Model

All code is in L2 memory (cached for cluster threads)




PULP Execution Model

Stacks are in L2 for the FC, and in **L1 (TCDM)** for the cluster cores (not cached!)



PULP Microcontroller Software Interface Standard (PMSIS)



PMSIS: How to Manage a Device Lifecycle

- **Configuration and initialization (*device specific*)**

`conf_init()`

`open_from_conf()`

- **Prepare the device for usage:**

`open()`

- **Perform required operations (*device specific*)**

- **Release the resources**

`close()`

PMSIS: Using the Cluster as a Device

```
struct pi_device cluster_dev = {0};  
struct pi_cluster_conf conf;  
struct pi_cluster_task cluster_task = {0};
```

} PMSIS data structures

```
// task allocation  
pi_cluster_task(&cluster_task, cluster_entry, NULL);
```

} Create a task for the cluster

```
// init the cluster  
pi_cluster_conf_init(&conf);  
pi_open_from_conf(&cluster_dev, &conf);
```

Function pointer

} Initialize the cluster device

```
// open the cluster  
if (pi_cluster_open(&cluster_dev)) return -1;
```

} Open the cluster device

```
// offload an entry point to the cluster  
pi_cluster_send_task_to_cl(&cluster_dev, &cluster_task);
```

} Execute code on the cluster

```
// releasing the cluster  
pi_cluster_close(&cluster_dev);
```

} Release the cluster device

Executing on the Cluster: Fork/join + SPMD

```
static void cluster_entry(void *arg)
```

```
{
```

```
    printf("Hello from cluster\n");
```

```
    // ...
```

```
    pi_cl_team_fork(NUM_CORES, cluster_fn, (void *) &args);
```

```
    // ...
```

```
}
```

} Executed on core 0

} Fork the execution of the same function on
NUM_CORES cores

} Executed on core 0

Data pointer

Function pointer

Synchronization: Barriers and Critical Sections

- **Barriers (*used for intermediate and final join points*):**


```
pi_cl_team_barrier();
```

- **Critical sections (*used to avoid critical races*):**

```
pi_cl_team_critical_enter();  
// code in the critical section  
pi_cl_team_critical_exit();
```

Hands-on Part, LAB 01 Recap:

Getting Started: *Docker setup*



- **How to install:**

- Install Docker Desktop* - <https://www.docker.com/products/docker-desktop/>
 - *For Windows users, it's recommended to use it through Windows Subsystem for Linux - WSL ([install instructions](#)); otherwise, it would become quite computationally-intensive, since it runs through a virtual machine.
- Install VS Code <https://code.visualstudio.com/>
- In VS Code, install the Dev Containers extension:
<https://marketplace.visualstudio.com/items?itemName=ms-vscode-remote.remote-containers>

Hands-on Part, LAB 01 Recap:

Getting Started: *Open project in the Docker container*

First steps:

1. Clone* this lab's project locally (in WSL, if using it)

```
$ git clone https://github.com/EEESlab/HSDDES-LAB04-PULP_Parallel_Bare.git
```

2. Enter the newly created local directory of the project

```
$ cd HSDDES-LAB04-PULP_Parallel_Bare
```

3. Open the project in VS Code

```
$ code .
```

At this moment, the project is open in your local “environment”.

We need to open it in the Docker container. For this, in VS Code:

1. **Press CTRL+Shift+P (open command palette)**
2. **Type in “Dev Containers: Reopen in Container” and press Enter**

We are now inside the container. From now on, when opening VS Code, the project will automatically open inside the Docker container.

*** We will also need to use the terminal from VS Code (CTRL+Shift+`) from now on, since it is the one accessing inside the running container!**



*In these labs we will always use Git. It is a popular tool, so a useful skill to have. Here is a good [intro tutorial](#).

Exercise 01: Vector Addition

- Now, take a look at the code and run the function on the cluster!

```
$ cd 01_vectAddPar  
$ make clean all run CORES=1
```

- How many clock cycles does it take SINGLE core ?

Ex. 01 - Method #1

01_vectAddPar/kernels.c

```
void vectAddPar(int * pSrcA, int * pSrcB, int * pDstC, int n) {  
    int i, core_id;
```

```
    core_id = pi_core_id();  
    for (i = core_id; i < n; i += NUM_CORES) {  
        pDstC[i] = a + b;  
    }
```

} Distribute single iterations among the available cores

```
    pi_cl_team_barrier();  
}
```

} Stop the execution once all cores have reached this point

- With NUM_CORES == 4 and n == 16:

• Iterations:	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
• Core id:	0	1	2	3	0	1	2	3	0	1	2	3	0	1	2	3

Ex. 01 - Method #2 - balanced workload

01_vectAddPar/kernels.c

```
void vectAddPar(int * pSrcA, int * pSrcB, int * pDstC, int n) {
    int i, core_id, i_start, i_end, i_block, i_chunk;
    core_id = pi_core_id();
    i_chunk = n / NUM_CORES;
    i_start = core_id * i_chunk;
    i_end = i_start + i_chunk;
    for (i = i_start; i < i_end; i++) {
        pDstC[i] = pSrcA[i] + pSrcB[i];
    }
    pi_cl_team_barrier();
}
```

} Split the workload into adjacent blocks (*chunks*) and compute their bounds [i_start, i_end)

- With NUM_CORES == 4 and n == 16:

• Iterations:	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
• Core id:	0	0	0	0	1	1	1	1	2	2	2	2	3	3	3	3

Ex. 01 - Method #2 - generic workload

01_vectAddPar/kernels.c

```
void vectAddPar(int * pSrcA, int * pSrcB, int * pDstC, int n) {  
    uint32_t i, core_id, i_chunk, i_start, i_end, i_block;  
    core_id = pi_core_id();  
    i_chunk = (n + NUM_CORES - 1) / NUM_CORES;  
    i_start = core_id * i_chunk;  
    i_end = i_start + i_chunk < n? i_start + i_chunk : n;  
  
    for (i = i_start; i < i_end; i++) {  
        pDstC[i] = pSrcA[i] + pSrcB[i];  
    }  
    pi_cl_team_barrier();  
}
```

1. Integer division with rounding towards infinity (1)
2. Check the upper bound of the iteration space

$$(1) \left\lceil \frac{a}{b} \right\rceil = \frac{a+b-1}{b}$$

- With NUM_CORES == 4 and n == 18:

• Iterations:	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17
• Core id:	0	0	0	0	0	1	1	1	1	1	2	2	2	2	2	3	3	3

Ex. 01 - Parallel Vector Add - Try on your own

- Apply the parallelization strategies by adapting the *vectAddPar* function.
- Run the code with different values of **CORES = {2, 4, 8}**
- Make sure the checksum output is correct.
- Measure the number of clock cycles and compute the parallel speedup= $\frac{time\ seq}{time\ par}$

```
$ make clean all run CORES=2  
$ make clean all run CORES=4  
$ make clean all run CORES=8
```

Exercise 02 - Parallelize a Matrix Multiplication Kernel

- This example is a “template” containing the code to run a matrix multiplication on the cluster. However, this is not optimized:

```
$ cd ../02_matrixMulPar  
$ make clean all run CORES=1
```

- Your task is to parallelize the code, using methods similar to the previous exercise.
- Measure the speed up and collect the performance statistics!
- Is the speed up ideal? If not, why?

Exercise 03: Parallelize an algorithm with data dependencies

1. Create a new application using the previous exercises as templates, and rename it “03_computePi”
2. Test the sequential code and print the result!

```
float iterative_pi(int n) {  
    float x, area=0.0f;  
    int i;  
    for(i=0; i<n; i++) {  
        x = (i + 0.5f)/n;  
        area += 4.0f/(1.0f + x*x);  
    }  
    return area/n;  
}
```

3. Parallelize the `iterative_pi` kernel.

**Hint: Make use of the critical section to update the `area` variable!*

*References on Makefile rules

- The Makefile is a “recipe” used to call the compiler/linker and, in this case, also to run the program. You can also chain several rules (e.g., `make clean all run`) and pass options (e.g., `runner_args="--vcd"`)
- Remove any existing build
`make clean`
- Build program (calling compiler + linker)
`make all`
- Run the program
`make run`
- Run options: you can change them by adding `runner_args="OPTIONS"`
`make run runner_args="--trace=.*insn.*" > trace.log` # written trace of instructions
- Disassembling:
Disassemble without inlined source code:
`make dis > test.S`
Disassemble with inlined source code:
`/pulp/pkg/pulp_riscv_gcc/bin/riscv32-unknown-elf-objdump -D -S > test.S`

Alberto Dequino – alberto.dequino@unibo.it

Călin Diaconu – calin.diaconu2@unibo.it

A graphic consisting of two overlapping white speech bubbles on a gray background. The top bubble is slightly offset to the left and contains the text "Q&A" in a bold, dark gray sans-serif font. The bottom bubble is slightly offset to the right and is empty.

Q&A

DEI – Università di Bologna

Viale del Risorgimento 2

Bologna, Italy



ALMA MATER STUDIORUM
UNIVERSITÀ DI BOLOGNA